

Deep Learning Review

machine learning

- conventional ml
 - were limited in their ability to process natural data in their raw form
 - required careful engineering and considerable domain expertise to design features
- representation learning
 - automatically discover the representations from raw data
 - deep learning
 - multiple levels of representation
 - transform the representation at one level into a representation at a higher, slightly more abstract level
 - good at discovering intricate structures in high-dimensional data
 - requires very little engineering by hand
 - take advantage of increases in the amount of available computation and data

supervised learning

- basic
 - labelled examples
 - adjustable weights
 - objective function that measures the error
- learning algorithm
 - computes a gradient vector
 - the negative gradient vector indicates the direction of steepest descent
 - stochastic gradient descent (SGD)
 - small set of examples
- generalization ability
 - its ability to produce sensible answers on new inputs that it has never seen during training
- linear classifiers
 - only carve their input space into very simple regions, namely half-spaces separated by a hyperplane
 - can not be insensitive to irrelevant variations of the input
 - hand design good feature extractors, which requires a considerable amount of engineering skill and domain expertise
- deep learning
 - good features can be learned automatically
 - multilayer stack of simple modules
 - can implement extremely intricate functions
 - sensitive to minute details
 - insensitive to large irrelevant variations

backpropagation

- train multilayer architectures
 - stochastic gradient descent
 - compute gradients using the backpropagation procedure
 - a practical application of the chain rule of derivatives
- late 1990s
 - forsaken by the machine-learning community and ignored by the computer-vision and speech-recognition communities
 - issues
 - poor local minima
 - saddle points
- revived around 2006
 - introduced unsupervised learning procedures that could create layers of feature detectors without requiring labelled data
- 2012
 - versions of the deep net from 2009 were being developed by many of the major speech groups and were already being deployed in Android phones

convolutional neural networks

- process data that come in the form of multiple arrays
- key ideas
 - local connections
 - shared weights
 - pooling
 - many layers
- stages
 - first
 - convolutional layers
 - to detect local conjunctions of features from the previous layer
 - pooling layers
 - to merge semantically similar features into one
 - second
 - two or three stages of convolution, non-linearity and pooling are stacked
 - followed by more convolutional and fully-connected layers

distributed representations and language processing

- two exponential advantages of deep nets that use distributed representations
 - enable generalization to new combinations of the values of learned features beyond those seen during training
 - composing layers of representation in a deep net brings the potential for another exponential advantage
- word vector
 - predict the next word in a sequence from a local context of earlier words
 - each word in the context is presented to the network as a one-of-N vector
 - in the first layer, each word creates a different pattern of activations, or word vectors

recurrent neural networks

- introduce
 - sequential inputs
 - one element at a time
 - state vector
 - historical information
- issue
 - gradients explode or vanish
 - grow or shrink at each time step
 - solved by advances in architecture and ways of training
 - predict next character or word
- usage
 - machine translation
 - encoder
 - generate state vector of the thought expressed by the sentence
 - decoder
 - state vector from encoder as input
 - outputs probability distribution of the words
 - image to word
 - encoder
 - CNN
 - decoder
 - RNN
- augment RNNs
 - conventional RNNs
 - it is difficult to learn to store information for very long
 - LSTM
 - memory cell
 - machine translation
 - Neural Turing Machine
 - 'tape-like' memory
 - read from or write to
 - be taught 'algorithms'
 - output a sorted list of symbol
 - memory networks
 - associative memory
 - question-answering
 - answering questions that require complex inference

reasoning and symbol manipulation

the future of deep learning

- unsupervised learning
- combining deep learning and reinforcement learning
- natural language understanding
- artificial intelligence come about